


Structure of Harmonia


- Structure of Harmonia consist of precisely two things.
 1. vector <struct Node> key_region
 2. vector <int> child_prefix_sum
- Apart from this two there one additional global variable used to initialize the child_prefix_sum
 1. psum



(b) Harmonia tree structure


Example

B+ tree :




fanout = 4

Key Region :



Child prefix sum :



Algorithm createHarmonia(Blocks , column_size)

Create a new vector of blocks called newBlocks

for each block inside Blocks do

 store block address of current Block node in curBlock named block pointer

 create temporary Node t and add keycount , all keys and all records in to it from the curBlock

 push the temporary Node t into the key_region vector

 push all the child of the curBlock into the newBlocks vector and increment psum for each child being pushed

 if it is first child of the curBlock then

 push psum to the child_prefix_sum vector

for end

If size of newBlocks vector is empty then

 clear Blocks vector and we are done.

else

 clear Blocks vector

 recursively call createHarmonia with newBlocks and column_size i.e. createHarmonia(newBlocks , column_size)

Searches

- Several other operations like Range queries and Update to the database requires the searching of the particular key firstly and then followed by that respective operation.
- Hence we will only discuss searches here in detail. Rest two operation are just variation of the search operation.
- Let's search the value 17 in our previously constructed structure of Harmonia.

Key Region :

7	14	21	24	5	7	9	11	14	17	20	21	24	26
0	1	2	3	4									

Child prefix sum :

1	4	4	4	4
---	---	---	---	---

So, How we will Search using given structure?

- Following equation will help in searching in current structure of Harmonia.

$$\text{child_idx} = \text{PrefixSum}[\text{node_idx}] + i - 1 \quad (1)$$


Searching : 17

$$(i) \text{child_id} = 1 + 3 - 1 = 3$$

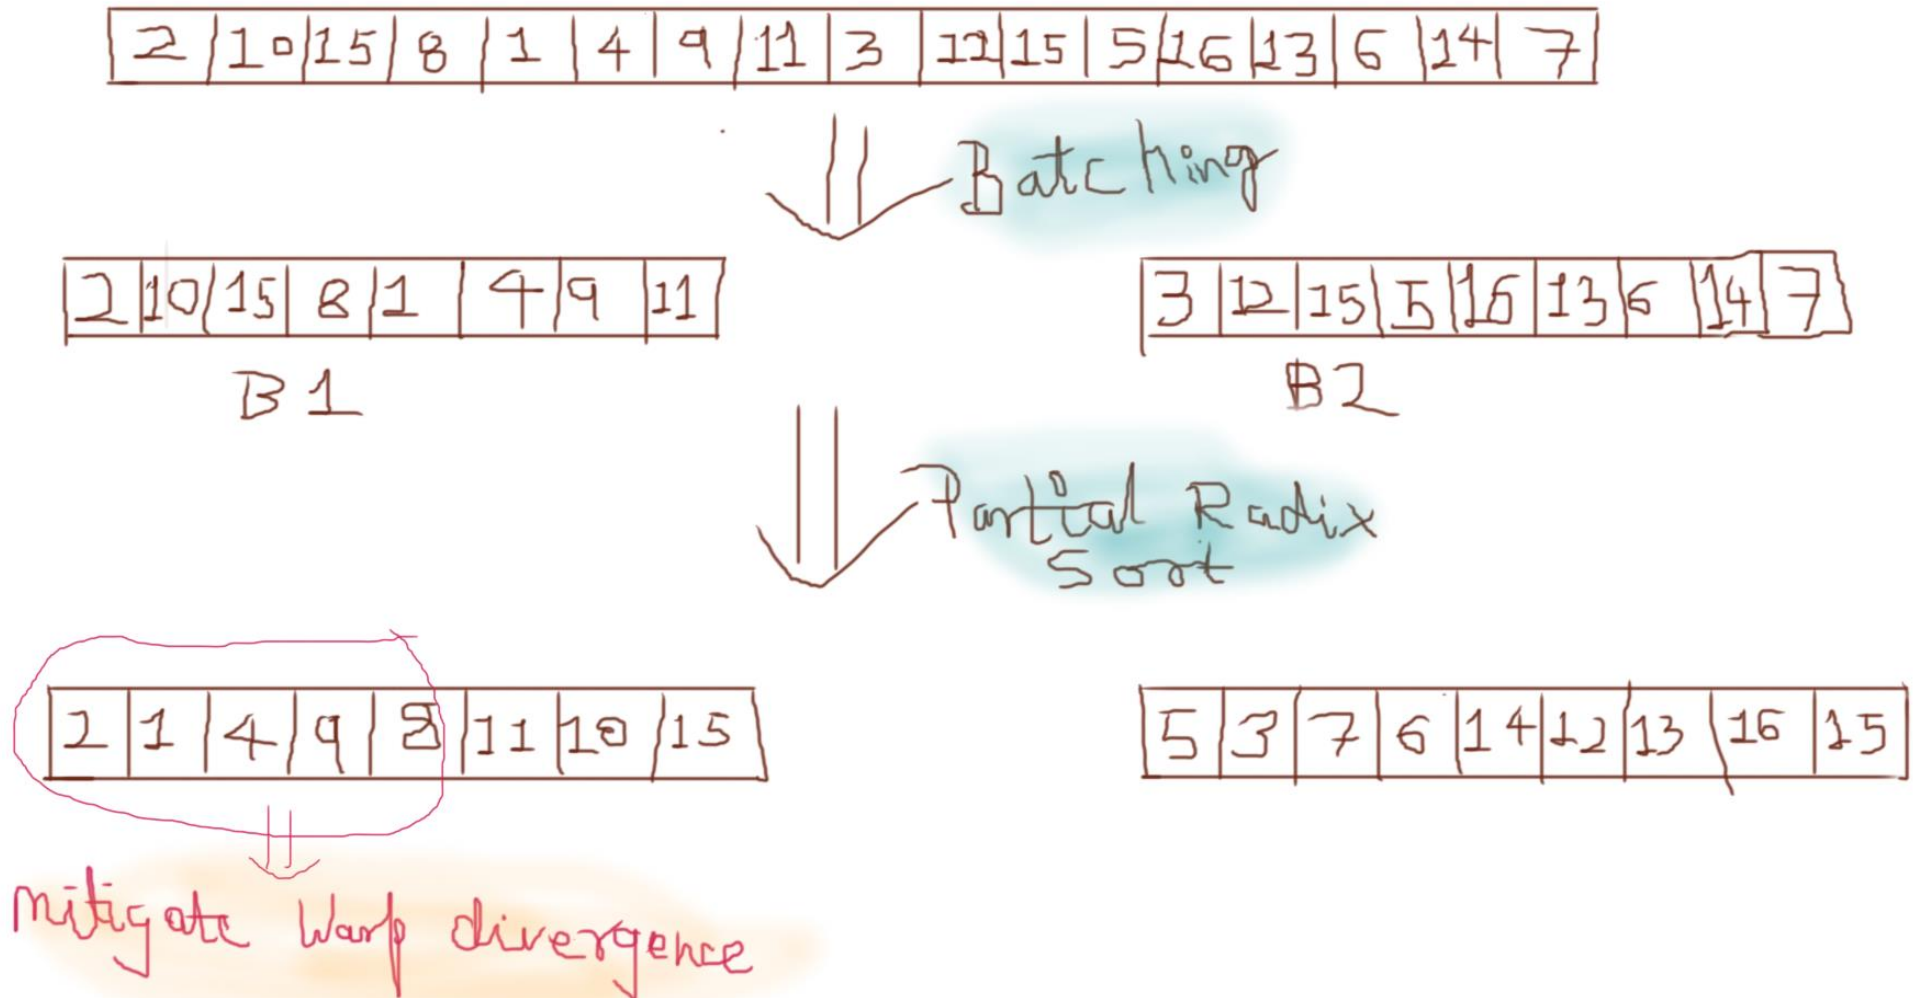
next level
node to be
Searched

Key Region :

child prefix sum :





Optimization 1: Batch-wise Partial Radix Sort on GPU (on search keys)



But what is the use of that?

- It makes memory access coalesced and also mitigates warp divergence



- Tuning of the number of MSB bits used for sorting is needed to get optimal performance.

$$N = B - \log_2\left(\frac{2^B}{T} * K\right)$$

B = each key is represented by B bits


T = the size of traversed B+ tree is T

K = cache line can save K keys

N = MSB N bits that has to be used for sorting by radix sort.

Optimization 2: Removing Unnecessary Computations (inside search procedure) or Narrow Thread Grouping

- Narrow the thread group based on the current number of keys present in the Node.



Procedure Search()


```
__global__ void search( struct Node *a , int *b , int asize , int bsize , int *search_keys, int *mutex , int n , char mode){  
    // task 1 is to assign individual searches to each threads  
    int key = search_keys[blockIdx.x];  
    __shared__ int index ;  
    __shared__ int prev_index;  
  
    index = 0;  
    prev_index = 0;  
  
    __syncthreads();  
  
    // task 2 is to perform search using harmonia  
    while(true){  
        if( index > asize-1 ){  
            break;  
        }  
        prev_index = index;  
        // need to divide this for loop among several threads to implement NTG.  
  
        if( threadIdx.x < a[index].count ){  
            if( a[index].keys[ threadIdx.x ][0] == key ){  
                index = b[index] + threadIdx.x + 1; ← eq 1  
                goto bottom;  
            }  
            if( threadIdx.x != a[index].count - 1 )  
            if( a[index].keys[ threadIdx.x ][0] < key  && key < a[index].keys[threadIdx.x + 1][0] ){  
                index = b[index] + threadIdx.x + 1;  
                goto bottom;  
            }  
        }  
        if( threadIdx.x == 0 ){  
            if( a[index].keys[0][0] > key ){  
                index = b[index] + 0;  
                goto bottom;  
            }  
            if( a[index].keys[ a[index].count-1 ][0] < key ){  
                index = b[index] + a[index].count;  
                goto bottom;  
            }  
        }  
    }  
    }  
    bottom: __syncthreads();  
}
```

How InsertKey works?


Case 1 :

Insert key
10

Non-Leaf



$10 < 15$
Value < curBlock \rightarrow value[i]
and
not null



Case 2 :
Insert key
10

$10 < \infty$
Value < curBlock \rightarrow value[i] and null

Algorithm InsertKey(curBlock, value , recordptr)

for each keys in the curBlock

 if value < curBlock->value[i] and curBlock->childBlock[i] is not empty then // case 1

 recursively call InsertKey on next level of tree pointed by that childBlock[i] i.e. InsertKey(curBlock->childBlock[i] ,
value , recordptr)

 if curBlock is at its max capacity

 split curBlock as Non Leaf node

 return as insertion is completed

 else if value < curBlock->value[i] and curBlock->childBlock[i] is empty then // case 2

 swap curBlock->value[i], and value

 update the number of keys present in curBlock



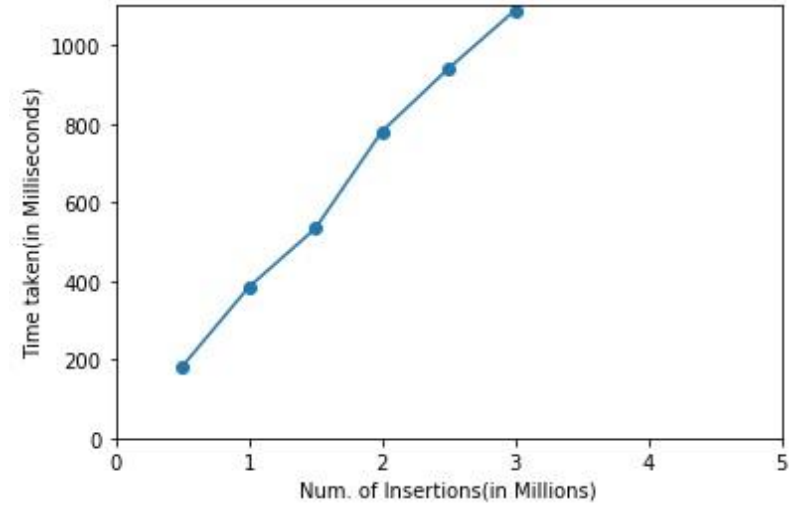

for end.

if curBlock is at its max capacity

 split curBlock as Leaf Node.

end.

but what about performance of all four operations?



And what about its comparison with CPU code?

- From the below Graph we can see that given implementation of Harmonia on GPU gives almost 23X speed up compared to the sequential code on CPU for Searches.

